

Report

Screenshot of running code

BCSF17A505

Mian Muhammad Anwar

Compiler Construction

CS 460: Cloud Computing

Punjab University College of Information Technology (PUCIT)

Assignment 3: Sentiment Analysis using Hadoop

Spring 2021

University of the Punjab

Instructor: Waheed Iqbal & Muhammad Abdullah

Assignment 03: Sentiment Analysis using Hadoop

Sentiment analysis is a process to identify emotions behind any text. In this assignment, you have to build a simple sentiment analysis MapReduce job. Following are the important files to perform the task:

1. twitter.json.tar.gz: A compressed file containing around 1 million tweets in a JSON format
2. positive_words.txt: A file containing list of positive words
3. negative_words.txt: A file containing list of negative words
4. stop_words.txt: A file containing list of stop words

Your MapReduce job should label each tweet either positive, negative, or neutral. For each tweet i , you need to ignore stop words, calculate positive words count p_i , and negative words count n_i . Then use the following formula to calculate the sentiment score s_i :

$$s_i = \frac{p_i - n_i}{p_i + n_i + 1} \quad (1)$$

A tweet i will be positive if s_i is greater than 0, negative if s_i is less than 0, and neutral if s_i is equal to 0. You must keep the original JSON document for each tweet and append additional attributed named *sentiment*. You need to produce three separate files for each category in JSON format.

Submission Guidelines

You need to prepare a small report contain screenshot of running your code, attach code in a zip file, and output file before **Tuesday 01 June, 11:59pm** and submit over Google Classroom.

Activities Terminal Wed 09:22

Twitts.java - Desktop - Visual Studio Code

EXPLORER

DESKTOP

03-Sentiment-Analysis

hadoop.java

negative_words.txt

positive_words.txt

stop_words.txt

Twitts.java U X

twitter.json

twitter.json (deleted)

03-Sentiment-Analysis

```
1 // Class-1 for Twitts
2 public class Twitts {
3     id_idObject;
4     private String user_id;
5     private String social_feed;
6     private String timestamp;
```

hadoop@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1

Shuffle Errors

BAD_ID=0

CONNECTION=0

IO_ERROR=0

WRONG_LENGTH=0

WRONG_MAP=0

WRONG_REDUCE=0

File Output Format Counters

Bytes Written=98

2021-06-16 09:21:16,822 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!

org.apache.hadoop.mapred.FileAlreadyExistsException: Output directory file:/home/hadoop/hadoop-3.2.1/tmp already exists

at org.apache.hadoop.mapreduce.lib.output.FileOutputFormat.checkOutputSpecs(FileOutputFormat.java:164)

at org.apache.hadoop.mapreduce.JobSubmitter.checkSpecs(JobSubmitter.java:277)

at org.apache.hadoop.mapreduce.JobSubmitter.submitJobInternal(JobSubmitter.java:143)

at org.apache.hadoop.mapreduce.Job511.run(Job.java:1570)

at org.apache.hadoop.mapreduce.Job511.run(Job.java:1567)

at java.security.AccessController.doPrivileged(Native Method)

at javax.security.auth.Subject.doAs(Subject.java:422)

at org.apache.hadoop.security.UserGroupInformation.doAs(UserGroupInformation.java:1730)

at org.apache.hadoop.mapreduce.Job.submit(Job.java:1567)

at org.apache.hadoop.mapreduce.Job.waitForCompletion(Job.java:1588)

at org.apache.hadoop.examples.Grep.run(Grep.java:94)

at org.apache.hadoop.util.ToolRunner.run(ToolRunner.java:76)

at org.apache.hadoop.examples.Grep.main(Grep.java:103)

at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)

at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)

at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

at java.lang.reflect.Method.invoke(Method.java:498)

at org.apache.hadoop.util.ProgramDriver\$ProgramDescription.invoke(ProgramDriver.java:71)

at org.apache.hadoop.util.ProgramDriver.run(ProgramDriver.java:144)

at org.apache.hadoop.examples.ExampleDriver.main(ExampleDriver.java:74)

at sun.reflect.NativeMethodAccessorImpl.invoke0(Native Method)

at sun.reflect.NativeMethodAccessorImpl.invoke(NativeMethodAccessorImpl.java:62)

at sun.reflect.DelegatingMethodAccessorImpl.invoke(DelegatingMethodAccessorImpl.java:43)

at java.lang.reflect.Method.invoke(Method.java:498)

at org.apache.hadoop.util.RunJar.run(RunJar.java:323)

at org.apache.hadoop.util.RunJar.main(RunJar.java:236)

> OUDhadoop@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1\$

40 public void set id id_idObject {

Ln 60, Col 32 Spaces: 4 UTF-8 LF Java

Activities Terminal Wed 09:22

Twitts.java - Desktop - Visual Studio Code

EXPLORER

DESKTOP

03-Sentiment-Analysis

hadoop.java

negative_words.txt

positive_words.txt

stop_words.txt

Twitts.java U X

twitter.json

twitter.json (deleted)

03-Sentiment-Analysis

```
1 // Class-1 for Twitts
2 public class Twitts {
3     id_idObject;
4     private String user_id;
5     private String social_feed;
6     private String timestamp;
```

hadoop@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1

2021-06-16 09:21:15,040 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2021-06-16 09:21:15,242 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties

2021-06-16 09:21:15,339 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).

2021-06-16 09:21:15,339 INFO Impl.MetricsSystemImpl: JobTracker metrics system started

2021-06-16 09:21:15,496 INFO Input.FileInputFormat: Total input files to process : 0

2021-06-16 09:21:15,517 INFO mapreduce.JobSubmitter: number of splits:0

2021-06-16 09:21:15,653 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_localS1404780_0001

2021-06-16 09:21:15,653 INFO mapreduce.JobSubmitter: Executing with tokens: []

2021-06-16 09:21:15,783 INFO mapreduce.Job: The url to track the job: http://localhost:8080/

2021-06-16 09:21:15,784 INFO mapreduce.Job: Running Job: job_localS1404780_0001

2021-06-16 09:21:15,780 INFO mapred.LocalJobRunner: OutputCommiter set to config null

2021-06-16 09:21:15,795 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2021-06-16 09:21:15,795 INFO output.FileOutputCommitter: skip cleanup temporary folders under output directory:false, ignore cleanup failures: false

2021-06-16 09:21:15,796 INFO mapred.LocalJobRunner: OutputCommiter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter

2021-06-16 09:21:15,833 INFO mapred.LocalJobRunner: map task executor complete.

2021-06-16 09:21:15,838 INFO mapred.LocalJobRunner: Waiting for reduce tasks

2021-06-16 09:21:15,838 INFO mapred.LocalJobRunner: Starting task: attempt_localS1404780_0001_r_0000000_0

2021-06-16 09:21:15,873 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2021-06-16 09:21:15,873 INFO output.FileOutputCommitter: skip cleanup temporary folders under output directory:false, ignore cleanup failures: false

2021-06-16 09:21:15,891 INFO mapred.Task: Using ResourceCalculatorProcessTree: []

2021-06-16 09:21:15,894 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle639ceb02

2021-06-16 09:21:15,896 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!

2021-06-16 09:21:15,915 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=3240257024, maxSingleShuffleLimit=810064256, mergeThreshold=2138569728, ioSortFactor=10, memToMemMergeOutputsThreshold=10

2021-06-16 09:21:15,917 INFO reduce.EventFetcher: attempt_localS1404780_0001_r_0000000_0 Thread started: EventFetcher for fetching Map Completion Events

2021-06-16 09:21:15,924 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning

2021-06-16 09:21:15,927 INFO mapred.LocalJobRunner:

2021-06-16 09:21:15,927 INFO reduce.MergeManagerImpl: finalMerge called with 0 in-memory map-outputs and 0 on-disk map-outputs

2021-06-16 09:21:15,928 INFO reduce.MergeManagerImpl: Merging 0 files, 0 bytes from disk

2021-06-16 09:21:15,929 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce

2021-06-16 09:21:15,932 INFO mapred.Merger: Merging 0 sorted segments

2021-06-16 09:21:15,932 INFO mapred.Merger: Down to the last merge-pass, with 0 segments left of total size: 0 bytes

2021-06-16 09:21:15,933 INFO mapred.LocalJobRunner:

2021-06-16 09:21:15,967 INFO Configuration.deprecation: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords

2021-06-16 09:21:15,979 INFO mapred.Task: Task:attempt_localS1404780_0001_r_0000000_0 is done. And is in the process of committing

2021-06-16 09:21:15,981 INFO mapred.LocalJobRunner:

> OUDhadoop@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1\$

40 public void set id id_idObject {

Ln 60, Col 32 Spaces: 4 UTF-8 LF Java

Activities Terminal Wed 09:22
Twitter.java - Desktop - Visual Studio Code

EXPLORER
03-Sentiment-Analysis
 hadoop.java
 negative_words.txt
 positive_words.txt
 stop_words.txt

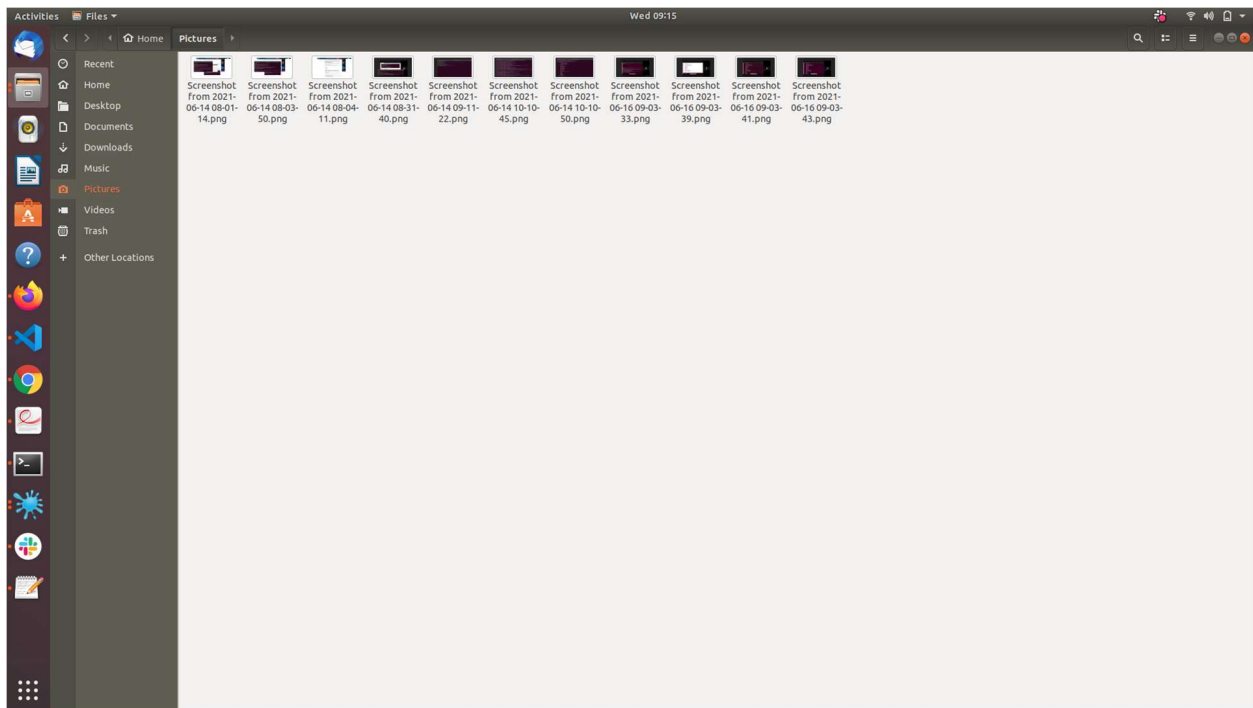
03-Sentiment-Analysis
1 // Class-1 for Twitts
2 public class Twitts {
3 id_idObject;
4 private String user_id;
5 private String social_feed;
6 private String timestamp;

hadoop@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1

```
2021-06-16 09:21:15.985 INFO mapred.LocalJobRunner: reduce > reduce
2021-06-16 09:21:15.985 INFO mapred.Task: Task 'attempt_local51404780_0001_r_000000_0' done.
2021-06-16 09:21:15.996 INFO mapred.Task: Final Counters for attempt_local51404780_0001_r_000000_0: Counters: 24
File System Counters
  FILE: Number of bytes read=316556
  FILE: Number of bytes written=835897
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=0
  Reduce shuffle bytes=0
  Reduce input records=0
  Reduce output records=0
  Spilled Records=0
  Shuffled Maps=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=8
  Total committed heap usage (bytes)=313524224
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=98
2021-06-16 09:21:15.996 INFO mapred.LocalJobRunner: Finishing task: attempt_local51404780_0001_r_000000_0
2021-06-16 09:21:15.997 INFO mapred.LocalJobRunner: reduce task executor complete.
2021-06-16 09:21:16.780 INFO mapreduce.Job: Job job_local51404780_0001 running in uber mode : false
2021-06-16 09:21:16.790 INFO mapreduce.Job: map 0% reduce 100%
2021-06-16 09:21:16.791 INFO mapreduce.Job: Job job_local51404780_0001 completed successfully
2021-06-16 09:21:16.798 INFO mapreduce.Job: Counters: 24
File System Counters
  FILE: Number of bytes read=316556
  FILE: Number of bytes written=835897
  FILE: Number of read operations=0
  FILE: Number of large read operations=0
  FILE: Number of write operations=0
Map-Reduce Framework
  Combine input records=0
  Combine output records=0
  Reduce input groups=0
  Reduce shuffle bytes=0
  Reduce input records=0
  Reduce output records=0
  Spilled Records=0
  Shuffled Maps=0
  Failed Shuffles=0
  Merged Map outputs=0
  GC time elapsed (ms)=8
  Total committed heap usage (bytes)=313524224
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
File Output Format Counters
  Bytes Written=98
```

master+ 40 public void set_id(id_idObject) {
Ln 60, Col 32 Spaces: 4 UTF-8 LF Java

```
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0
Job Counters
  Launched map tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=2238
  Total time spent by all reduces in occupied slots (ms)=2584
  Total time spent by all map tasks (ms)=2239
  Total time spent by all reduce tasks (ms)=2584
  Total vcore-milliseconds taken by all map tasks=2239
  Total vcore-milliseconds taken by all reduce tasks=2584
  Total megabyte-milliseconds taken by all map tasks=2292736
  Total megabyte-milliseconds taken by all reduce tasks=2564896
Map-Reduce Framework
  Map input records=3
  Map output records=0
  Map output bytes=88
  Map output materialized bytes=80
  Input split bytes=110
  Combine input records=0
  Combine output records=0
  Reduce input groups=0
  Reduce shuffle bytes=80
  Reduce input records=0
  Reduce output records=0
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=154
  CPU time spent (ms)=1450
  Physical memory (bytes) snapshot=524750848
  Virtual memory (bytes) snapshot=5219487744
  Total committed heap usage (bytes)=458130560
  Peak Map Physical memory (bytes)=292489344
  Peak Map Virtual memory (bytes)=2086050048
  Peak Reduce Physical memory (bytes)=232341304
  Peak Reduce Virtual memory (bytes)=2012637096
Shuffle Errors
  BAD_ID=0
  CONNECTION=0
  IO_ERROR=0
  WRONG_LENGTH=0
  WRONG_MAP=0
  WRONG_REDUCE=0
```




```
Activities Terminal + 2021-05-24 Mon 10:10
avg-qyxc-qj (2021-05-24) x
drive.google.com/file/d/1KAqPbNwOnaPkbFwLwK9XNq4QcNrxI/view

hdoo@dev-HP-ProBook-450-G7: ~/hadoop-3.2.1
File Edit View Search Terminal Help
hdoo@dev-HP-ProBook-450-G7:~/hadoop-3.2.1$ bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar grep input output 'dfs[a-z.]'
2021-06-14 10:10:33.067 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
2021-06-14 10:10:33.222 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties
2021-06-14 10:10:33.286 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).
2021-06-14 10:10:33.286 INFO Impl.MetricsSystemImpl: JobTracker metrics system started
2021-06-14 10:10:33.379 INFO Input.FileInputFormat: Total input files to process : 0
2021-06-14 10:10:33.392 INFO mapreduce.JobSubmitter: number of splits:0
2021-06-14 10:10:33.484 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1671980215_0001
2021-06-14 10:10:33.484 INFO mapreduce.JobSubmitter: Executing with tokens: []
2021-06-14 10:10:33.585 INFO mapreduce.Job: The url to track the job: http://localhost:8080/
2021-06-14 10:10:33.585 INFO mapreduce.Job: Running Job: job_local1671980215_0001
2021-06-14 10:10:33.586 INFO mapred.LocalJobRunner: OutputCommitter set in config null
2021-06-14 10:10:33.592 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-06-14 10:10:33.592 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2021-06-14 10:10:33.593 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter
2021-06-14 10:10:33.618 INFO mapred.LocalJobRunner: Waiting for map tasks
2021-06-14 10:10:33.618 INFO mapred.LocalJobRunner: map task executor complete.
2021-06-14 10:10:33.621 INFO mapred.LocalJobRunner: Waiting for reduce tasks
2021-06-14 10:10:33.622 INFO mapred.LocalJobRunner: Starting task: attempt_local1671980215_0001_r_000000_0
2021-06-14 10:10:33.640 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2
2021-06-14 10:10:33.640 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false
2021-06-14 10:10:33.652 INFO mapred.Task: Using ResourceCalculatorProcessTree : [ ]
2021-06-14 10:10:33.654 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@343f0d
2021-06-14 10:10:33.655 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!
2021-06-14 10:10:33.660 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=3240257024, maxSingleShuffleLimit=81064256, mergeThreshold=2138569728, ioSortFactor=10, memTokenMergeOutputsThreshold=10
2021-06-14 10:10:33.668 INFO reduce.EventFetcher: attempt_local1671980215_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events
2021-06-14 10:10:33.672 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning
2021-06-14 10:10:33.676 INFO reduce.MergeManagerImpl: finalMerge called with 0 in-memory map-outputs and 0 on-disk map-outputs
2021-06-14 10:10:33.677 INFO reduce.MergeManagerImpl: Merging 0 files, 0 bytes from disk
2021-06-14 10:10:33.677 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce
2021-06-14 10:10:33.680 INFO mapred.Merger: Merging 0 sorted segments
2021-06-14 10:10:33.680 INFO mapred.Merger: Down to the last merge-pass, with 0 segments left of total size: 0 bytes
2021-06-14 10:10:33.681 INFO mapred.LocalJobRunner:
2021-06-14 10:10:33.677 INFO mapred.ReduceTask: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords
2021-06-14 10:10:33.718 INFO mapred.Task: Task:attempt_local1671980215_0001_r_000000_0 is done. And is in the process of committing
2021-06-14 10:10:33.719 INFO mapred.LocalJobRunner:
2021-06-14 10:10:33.719 INFO mapred.Task: Task attempt_local1671980215_0001_r_000000_0 is allowed to commit now
2021-06-14 10:10:33.721 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1671980215_0001_r_000000_0' to file:/home/hdoo/hadoop-3.2.1/grep-temp-1466437422
2021-06-14 10:10:33.721 INFO mapred.LocalJobRunner: reduce > reduce
2021-06-14 10:10:33.721 INFO mapred.Task: Task 'attempt_local1671980215_0001_r_000000_0' done.
```

Activities Terminal + 2021-05-24 Mon 08:31

avg-qyxc-qj (2021-05-24) x

drive.google.com/file/d/1KAqPbNwOnaPkbFwLwK9XNq4QcNrxI/view

Apache Hadoop 3.2.1

bin/hadoop jar share/hadoop/mapreduce/hadoop-mapreduce-examples-3.2.1.jar grep input output 'dfs[a-z.]'

2021-06-14 10:10:33.067 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable

2021-06-14 10:10:33.222 INFO Impl.MetricsConfig: Loaded properties from hadoop-metrics2.properties

2021-06-14 10:10:33.286 INFO Impl.MetricsSystemImpl: Scheduled Metric snapshot period at 10 second(s).

2021-06-14 10:10:33.286 INFO Impl.MetricsSystemImpl: JobTracker metrics system started

2021-06-14 10:10:33.379 INFO Input.FileInputFormat: Total input files to process : 0

2021-06-14 10:10:33.392 INFO mapreduce.JobSubmitter: number of splits:0

2021-06-14 10:10:33.484 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_local1671980215_0001

2021-06-14 10:10:33.484 INFO mapreduce.JobSubmitter: Executing with tokens: []

2021-06-14 10:10:33.585 INFO mapreduce.Job: The url to track the job: http://localhost:8080/

2021-06-14 10:10:33.585 INFO mapreduce.Job: Running Job: job_local1671980215_0001

2021-06-14 10:10:33.586 INFO mapred.LocalJobRunner: OutputCommitter set in config null

2021-06-14 10:10:33.592 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2021-06-14 10:10:33.592 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false

2021-06-14 10:10:33.593 INFO mapred.LocalJobRunner: OutputCommitter is org.apache.hadoop.mapreduce.lib.output.FileOutputCommitter

2021-06-14 10:10:33.618 INFO mapred.LocalJobRunner: Waiting for map tasks

2021-06-14 10:10:33.618 INFO mapred.LocalJobRunner: map task executor complete.

2021-06-14 10:10:33.621 INFO mapred.LocalJobRunner: Waiting for reduce tasks

2021-06-14 10:10:33.622 INFO mapred.LocalJobRunner: Starting task: attempt_local1671980215_0001_r_000000_0

2021-06-14 10:10:33.640 INFO output.FileOutputCommitter: File Output Committer Algorithm version is 2

2021-06-14 10:10:33.640 INFO output.FileOutputCommitter: FileOutputCommitter skip cleanup _temporary folders under output directory:false, ignore cleanup failures: false

2021-06-14 10:10:33.652 INFO mapred.Task: Using ResourceCalculatorProcessTree : []

2021-06-14 10:10:33.654 INFO mapred.ReduceTask: Using ShuffleConsumerPlugin: org.apache.hadoop.mapreduce.task.reduce.Shuffle@343f0d

2021-06-14 10:10:33.655 WARN Impl.MetricsSystemImpl: JobTracker metrics system already initialized!

2021-06-14 10:10:33.660 INFO reduce.MergeManagerImpl: MergeManager: memoryLimit=3240257024, maxSingleShuffleLimit=81064256, mergeThreshold=2138569728, ioSortFactor=10, memTokenMergeOutputsThreshold=10

2021-06-14 10:10:33.668 INFO reduce.EventFetcher: attempt_local1671980215_0001_r_000000_0 Thread started: EventFetcher for fetching Map Completion Events

2021-06-14 10:10:33.672 INFO reduce.EventFetcher: EventFetcher is interrupted.. Returning

2021-06-14 10:10:33.676 INFO reduce.MergeManagerImpl: finalMerge called with 0 in-memory map-outputs and 0 on-disk map-outputs

2021-06-14 10:10:33.677 INFO reduce.MergeManagerImpl: Merging 0 files, 0 bytes from disk

2021-06-14 10:10:33.677 INFO reduce.MergeManagerImpl: Merging 0 segments, 0 bytes from memory into reduce

2021-06-14 10:10:33.680 INFO mapred.Merger: Merging 0 sorted segments

2021-06-14 10:10:33.680 INFO mapred.Merger: Down to the last merge-pass, with 0 segments left of total size: 0 bytes

2021-06-14 10:10:33.681 INFO mapred.LocalJobRunner:

2021-06-14 10:10:33.677 INFO mapred.ReduceTask: mapred.skip.on is deprecated. Instead, use mapreduce.job.skiprecords

2021-06-14 10:10:33.718 INFO mapred.Task: Task:attempt_local1671980215_0001_r_000000_0 is done. And is in the process of committing

2021-06-14 10:10:33.719 INFO mapred.LocalJobRunner:

2021-06-14 10:10:33.719 INFO mapred.Task: Task attempt_local1671980215_0001_r_000000_0 is allowed to commit now

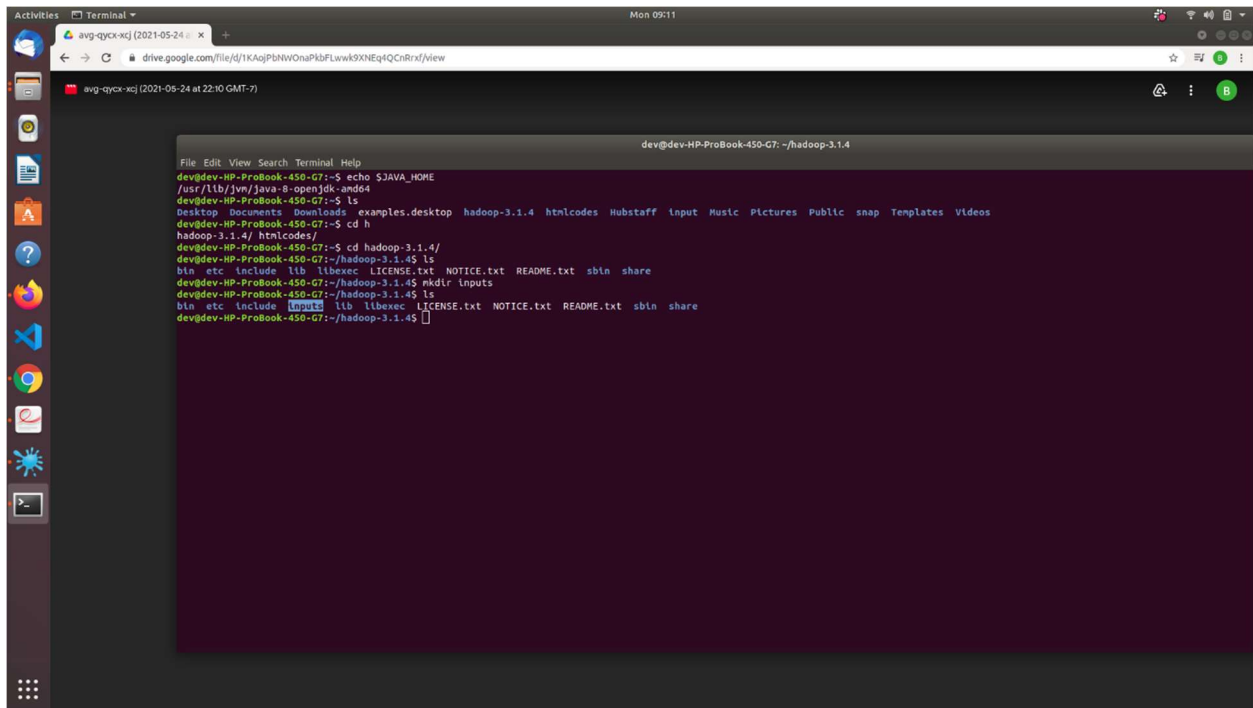
2021-06-14 10:10:33.721 INFO output.FileOutputCommitter: Saved output of task 'attempt_local1671980215_0001_r_000000_0' to file:/home/hdoo/hadoop-3.2.1/grep-temp-1466437422

2021-06-14 10:10:33.721 INFO mapred.LocalJobRunner: reduce > reduce

2021-06-14 10:10:33.721 INFO mapred.Task: Task 'attempt_local1671980215_0001_r_000000_0' done.

hdoo@dev-HP-ProBook-450-G7: ~

```
File Edit View Search Terminal Help
hdoo@dev-HP-ProBook-450-G7:~$ which javac
/usr/bin/javac
hdoo@dev-HP-ProBook-450-G7:~$ which java
/usr/bin/java
hdoo@dev-HP-ProBook-450-G7:~$ ls -ltr /usr/bin/java
lrwxrwxrwx 1 root root 22 Jun 14 07:49 /usr/bin/java -> /etc/alternatives/java
hdoo@dev-HP-ProBook-450-G7:~$ ls -ltr /etc/alternatives/java
lrwxrwxrwx 1 root root 46 Jun 14 07:49 /etc/alternatives/java -> /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java
hdoo@dev-HP-ProBook-450-G7:~$ ls -ltr /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java
-rwxr-xr-x 1 root root 6288 Apr 21 07:46 /usr/lib/jvm/java-8-openjdk-amd64/jre/bin/java
hdoo@dev-HP-ProBook-450-G7:~$
```



Activities Firefox Web Browser Mon 08:04

Cloud-Computing-Proje... Google Cloud Computing BS CS How to Install Hadoop on Apache Download Mirr...

https://phoenixnap.com/kb/install-hadoop-ubuntu

Bare Metal Cloud now available at special prices! SEE DISCOUNTS

The Hadoop binary files are now located within the `hadoop-3.2.1` directory.

Single Node Hadoop Deployment (Pseudo-Distributed Mode)

Hadoop excels when deployed in a **fully distributed mode** on a large cluster of networked servers. However, if you are new to Hadoop and want to explore basic commands or test applications, you can configure Hadoop on a single node.

This setup, also called **pseudo-distributed mode**, allows each Hadoop daemon to run as a single Java process. A Hadoop environment is configured by editing a set of configuration files:

- `bashrc`
- `hadoop-env.sh`
- `core-site.xml`
- `hdfs-site.xml`
- `mapred-site.xml`
- `yarn-site.xml`

Configure Hadoop Environment Variables (bashrc)

Edit the `bashrc` shell configuration file using a text editor of your choice (we will be using nano):

```
sudo nano .bashrc
```

Define the Hadoop environment variables by adding the following content to the end of the file:

```
#Hadoop Related Options
export HADOOP_HOME=/home/hadoop/hadoop-3.2.1
export HADOOP_INSTALL=$HADOOP_HOME
export HADOOP_MAPRED_HOME=$HADOOP_HOME
export HADOOP_COMMON_HOME=$HADOOP_HOME
export HADOOP_HDFS_HOME=$HADOOP_HOME
export YARN_HOME=$HADOOP_HOME
export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native
```

```
HDFS: Number of write operations=2
HDFS: Number of bytes read erasure-coded=0

Job Counters
  Launched map tasks=1
  Launched reduce tasks=1
  Data-local map tasks=1
  Total time spent by all maps in occupied slots (ms)=2239
  Total time spent by all reduces in occupied slots (ms)=2504
  Total time spent by all map tasks (ms)=2239
  Total time spent by all reduce tasks (ms)=2504
  Total vcore-milliseonds taken by all map tasks=2239
  Total vcore-milliseonds taken by all reduce tasks=2504
  Total megabyte-milliseonds taken by all map tasks=2292736
  Total megabyte-milliseonds taken by all reduce tasks=2564096

Map-Reduce Framework
  Map input records=3
  Map output records=8
  Map output bytes=83
  Map output materialized bytes=80
  Input split bytes=110
  Combine input records=8
  Combine output records=6
  Reduce input groups=6
  Reduce shuffle bytes=80
  Reduce input records=6
  Reduce output records=6
  Spilled Records=12
  Shuffled Maps =1
  Failed Shuffles=0
  Merged Map outputs=1
  GC time elapsed (ms)=154
  CPU time spent (ms)=1450
  Physical memory (bytes) snapshot=524750048
  Virtual memory (bytes) snapshot=5219487744
  Total committed heap usage (bytes)=456130560
  Peak Map Physical memory (bytes)=292409344
  Peak Map Virtual memory (bytes)=2006850048
  Peak Reduce Physical memory (bytes)=232341504
  Peak Reduce Virtual memory (bytes)=2612637696

Shuffle
  Errors
    BAD_ID=0
    CONNECTION=0
    IO_ERROR=0
    WRONG_LENGTH=0
    WRONG_MAP=0
    WRONG_REDUCE=0

File Input Format Counters
  Bytes Read=51
```

